

Bruno

API client with file-based collections and Git-friendly workflows.

What is it?

Bruno is an API client focused on local, versionable collections that work well with Git.

Official site and documentation

Use the official Bruno source as the primary reference: <https://www.usebruno.com/downloads>. Official documentation should be the first source for installation, supported versions, syntax changes and new features.

Basic use

Create a collection, configure method, URL, headers, parameters, body and authentication, then send the request.

Version control

File-based collections can be reviewed through diffs, commits and pull requests.

Secrets

Keep sensitive values out of commits and use appropriate variable/secret mechanisms.

When to consider

Useful for teams that want API testing close to a Git workflow.

When to use it

Use this knowledge when product risk, architecture or the delivery stage requires this type of validation. Confirm the goal, environment, required data and expected evidence before starting.

How to start safely

Start from the official Bruno site and documentation. Confirm prerequisites, supported versions, operating system and required permissions. Learn in a controlled environment before using it in a critical workflow.

First QA-oriented use

Pick one small, verifiable task in Bruno, configure only what is necessary, run a simple example and confirm the outcome. Repeat with different data or conditions before adding abstractions.

Using it in a real project

Define where Bruno fits, who owns configuration, how credentials and data are handled and how results are shared. Integrate it into the team workflow instead of creating an isolated activity.

Adoption good practices

Standardize versions and configuration, document important decisions, keep secrets out of code and preserve reproducible examples. Reassess whether Bruno still delivers value as the product evolves.

Common issues and troubleshooting

When Bruno fails, verify versions, dependencies, permissions, connectivity, data and environment before assuming a product defect. Preserve the original error and reduce the issue to the smallest reproducible case.